Niranjan Khedkar

Bengaluru, India | +91 8928571478 | nrnjnkhdkr@gmail.com | linkedin.com/in/niranjan-khedkar123 | niranjankhedkar.github.io

Summary

AI Engineer with 3 years at IBM building and operating production GenAI systems for global enterprise clients. Hands-on across the full lifecycle: LLM fine-tuning, RAG and agentic pipelines, GPU inference optimization, containerized deployment on Kubernetes/OpenShift, and multi-cloud production operations. Delivered 15+ production AI systems generating $10M+ in revenue for Japanese and Fortune 500 clients.

Skills

Languages: Python, SQL, Golang, JavaScript, Bash

GenAI & LLMs: LLM Fine-tuning (LoRA, SFT), RAG, Agentic AI (CrewAI, LangGraph), LangChain, LlamaIndex, MCP, vLLM, Hugging Face, Prompt Engineering, Vector Databases (Milvus, ChromaDB, FAISS), LLM Evaluation (RAGAS, BLEU, ROUGE, COMET, BLEURT)

Machine Learning: PyTorch, TensorFlow, Scikit-learn, XGBoost, Transformers (BERT), YOLO (Object Detection), OpenCV, Multi-modal AI (Image-to-Text)

MLOps & Cloud: Docker, Kubernetes, OpenShift, Helm, AWS, GCP, Azure, CI/CD, Monitoring & Inference Optimization

Data & Backend: Pandas, NumPy, PostgreSQL, FastAPI, Flask, Django, REST APIs, Streamlit

Spoken Languages: English (Native), Japanese (JLPT N4, working proficiency with clients)

Experience

IBM Bengaluru, India
AI Delivery Consultant | Applied AI *Nov 2025 – Present*

- Architect and operate production AI/ML workloads on OpenShift and Kubernetes, owning orchestration, autoscaling, and resource optimization for containerized inference services.
- Operate enterprise AI systems for Fortune 500 clients across logistics, enterprise software, and telecommunications, maintaining **99.9%+ uptime** for mission-critical automation workloads.
- Fine-tune deployed models and automation workflows against production feedback, improving inference speed, accuracy, and resource consumption per client SLAs.
- Troubleshoot complex production incidents across multi-cloud environments (AWS, GCP, Azure, IBM Cloud), reducing mean time to resolution through systematic debugging and runbook frameworks.

IBM Bengaluru, India
AI Engineer | Forward Deployed Engineer (FDE) *Oct 2023 – Oct 2025*

- Architected and deployed **15+ production GenAI systems worth $10M+ in revenue** for Japanese enterprise clients (leading automotive manufacturers and financial institutions), using enterprise LLMs with RAG pipelines and autonomous agents.
- Fine-tuned **IBM Granite and Llama** models on domain-specific datasets (**100K+ documents**), improving response accuracy and contextual relevance by **20–30%** in production.
- Cut manual research and coding time by up to **60%** via multi-agent systems (CrewAI, LangGraph) orchestrating retrieval, summarization, reasoning, and tool invocation across APIs, databases, and enterprise knowledge bases.
- Delivered AI-driven code generation and conversion pipelines for a leading global automotive manufacturer and a major consumer-electronics company, accelerating software migration and boosting developer productivity by **40%**.
- Built hybrid-cloud AI infrastructure (Docker, OpenShift) with end-to-end MLOps and rigorous LLM evaluation (perplexity, BLEU, ROUGE, human-in-the-loop) plus automated feedback loops.
- Built computer vision and multi-modal AI solutions, including image-to-text generation and YOLO-based symbol detection and classification for nuclear plant engineering diagrams.

Richin AI Mumbai, India (Remote)
AI Engineer (Freelance) | Python, NLP, Chatbots *Jul 2023 – Aug 2023*

- Designed and deployed a customer-facing chatbot for a finance company, cutting response time by **40%** and improving sentiment analysis accuracy by **18%**.

Projects

Neural Machine Translation Platform – ESA Partner | *vLLM, NVIDIA A100, Gemma 2/3, LoRA, SFT* May 2025 – Nov 2025

- Primary technical resource for a global language-services provider's neural translation platform; migrated the serving stack from AWS to cloud GPU infrastructure (2× NVIDIA A100 80GB) via a BYOM vLLM runtime with continuous batching.
- Designed an asynchronous proxy layer that load-balanced and dispatched batch translation requests in parallel across two vLLM model replicas, achieving **6.2× inference throughput** (2712s → 280s for 2,000 segments).
- Built an end-to-end fine-tuning pipeline with Hugging Face: corpus prep, base model benchmarking (Gemma 2 9B / Gemma 3 12B/27B), LoRA + SFT continual pre-training, and multi-metric evaluation (BLEU, COMET, BLEURT, MetricX-24 XL); delivery across three time zones **secured the ESA deal**.

Project Hail Mary – Offline Audio Translator | *Python, Librosa, FAISS, ONNX Runtime, Streamlit* Personal

- Built a fully offline sound-to-word translator inspired by Andy Weir's novel: extracts 40 MFCC features per audio clip via Librosa, indexes them in FAISS, and retrieves the closest match by L2 distance.

Education

Pillai College of Engineering, University of Mumbai Mumbai, India
Bachelor of Technology (CGPA: 9.61/10) *Jul 2019 – Jul 2023*

Publications, Achievements & Certifications

Publication: *Determination of Medicinal Leaf Properties*, IJSREM – CNN-based plant classification research (94% accuracy)

IBM Employee Equity Program (EEP) Award | **IBM 100% Club Member** | **JLPT N4 Certified** (Japanese)